

AER1517 Control for Robotics


Assignment 2: Markov Decision Processes, Classical Reinforcement Learning & Model Predictive Control

Prof. Angela Schoellig

- Download handout and script templates from Quercus
 - Problem 2.1 Grid World
 - Problem 2.2 Mountain Car
- Due on Mar. 24 (Tuesday) 23:59
- Submission through *Gradescope*
 - A single PDF with solutions and requested scripts
 - Both typed and scanned handwritten solutions are accepted
- Office hour: Mar. 19 (Thursday) 14:00 via Hangout or in-person @UTIAS
- Submit email questions by Mar. 15 or Mar. 22 (Sunday) --- open issues are discussed during lectures


Problem 1.1 Grid World | Overview

- Goal: Plan a path from 'S' to 'G' without transitioning to obstacle states



Problem 1.1 Grid World | Overview

- Goal: Plan a path from 'S' to 'G' without transitioning to obstacle states



- Task: Solve the grid world problem with
 - Model-based approach: Generalized Policy Iteration (GPI)
 - Sample-based approaches: Monte-Carlo (on-policy) and Q Learning (off-policy)
 - **[Extra Credit]** Linear Programming

Problem 1.1 Grid World | Code Structure

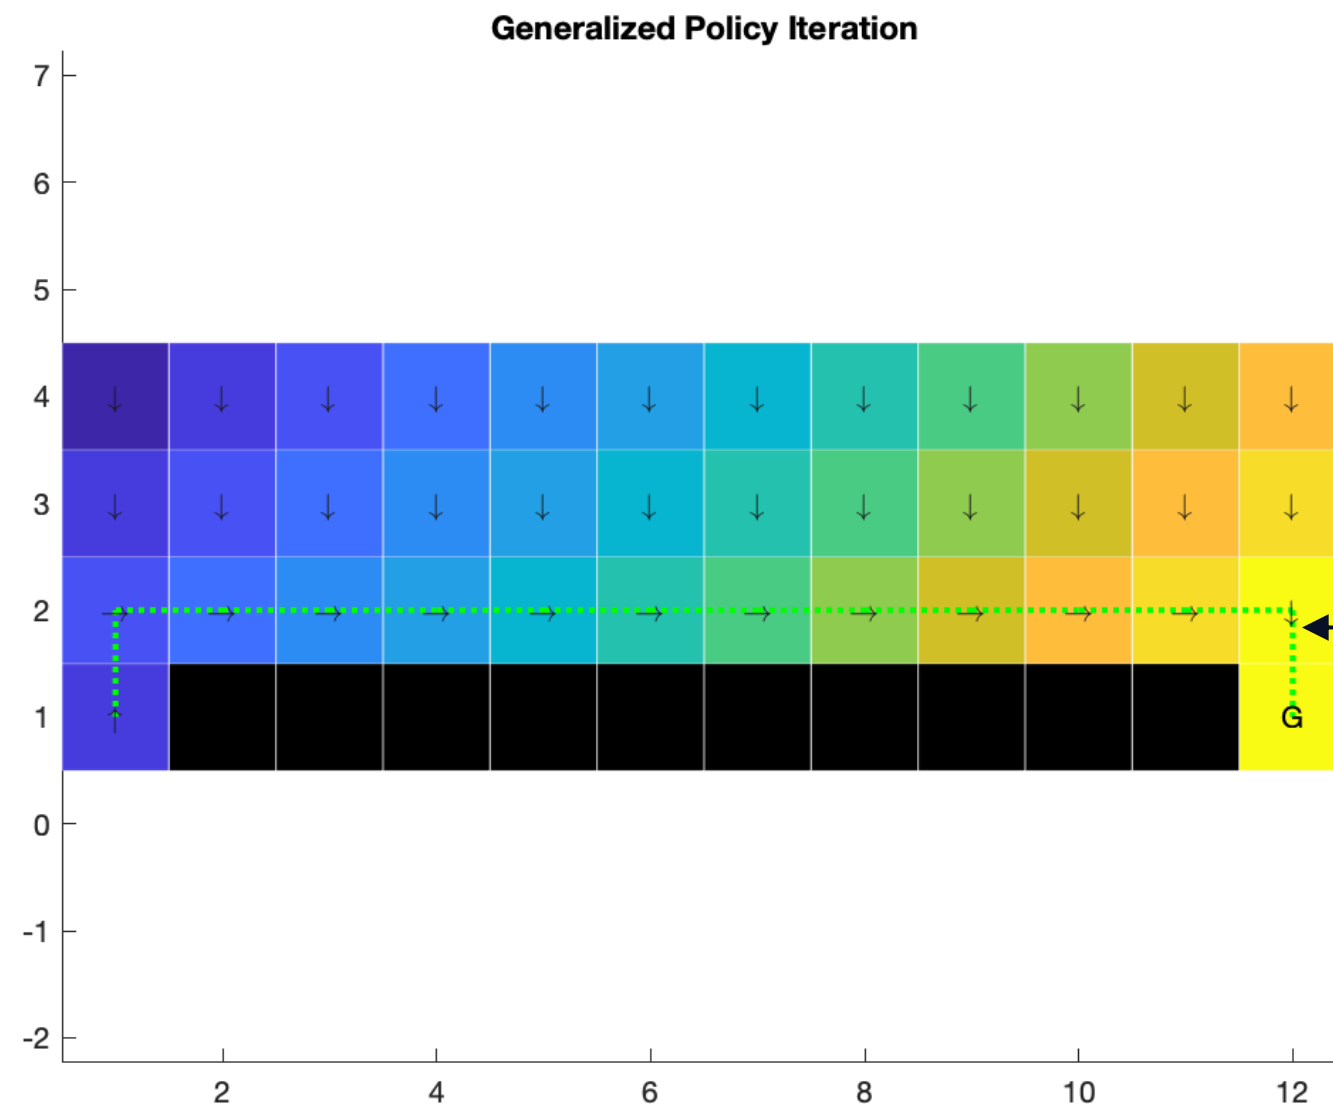
main_p1_gw.m

- Load model
- References
- Default parameters
- Algorithm (e.g., GPI)
- Visualization



```
EDITOR PUBLISH VIEW
reinterp_init_impact.m netExtractParameters2.m main_p1_gw.m +
42 %% General
43 % Load world
44 load('./gridworld_model/grid_world');
45
46 % Add path
47 addpath(genpath(pwd));
48
49 % Result and plot directory
50 save_dir = './results/';
51 mkdir(save_dir);
52
53 %% Problem 2.1: (a) Generalized Policy Iteration (GPI)
54 % Instruction: Implement the GPI algorithm to solve the grid world problem
55 % Reference: Section 2.8 of [1]
56
57 % Parameters of the GPI algorithm
58 precision_pi = 0.1;
59 precision_pe = 0.01;
60 max_ite_pi = 100;
61 max_ite_pe = 100;
62
63 % ===== [TODO] GPI Implementation =====
64 % Complete implementation in 'generalized_policy_iteration'
65 [v_gpi, policy_gpi] = generalized_policy_iteration(world, precision_pi, ...
66     precision_pe, max_ite_pi, max_ite_pe);
67 % =====
68
69 % Visualization
70 plt_title = 'Generalized Policy Iteration';
71 plt_path = true;
72 plt_gpi = visualize_gw_solution(world, v_gpi, policy_gpi, ...
73     plt_title, plt_path);
74
75 % Save results and figure to report
76 save(strcat(save_dir, 'gpi_results.mat'), 'v_gpi', 'policy_gpi');
77 saveas(plt_gpi, strcat(save_dir, 'gpi_plot.png'), 'png');
```

Problem 1.1 Grid World | Example Result with GPI



Heatmap corresponds to state value function (or, cost-to-go)

Planned path from the start state to the goal state

Problem 2.2 Mountain Car | Overview

- Goal: Drive an under-powered car to the top of the hill

- States: position x and velocity v

- Input: acceleration a

- System dynamics:

$$v_{k+1} = v_k + 0.001a_k - 0.0025 \cos(3x_k)$$


$$x_{k+1} = x_k + v_{k+1}$$

- Bounds on states and input:

$$x_k \in [-1.2, 0.5], v_k \in [-0.07, 0.07], a_k \in [-1, 1]$$

- Task:

- Create stochastic Markov Decision Process (MDP) and solve with classical RL
- Formulate and solve with Model Predictive Control (MPC)



Problem 1.2 Mountain Car | Code Structure (MDP Design)

create_mountain_car.m

- Build MDP



```
EDITOR PUBLISH VIEW
create_mountain_car.m x main_p2_mc_rl.m x +
106 %% Problem 2.2: (a)-(b) Create stochastic MDPs for the mountain car problem
107 % Instruction: Implement the Nearest Neighbour and Linear Interpolation
108 % approaches for creating discrete stochastic MDPs
109 % Reference: see [4] for linear_interp implementation
110
111 % ===== [TODO] Discretization =====
112 % Complete implementations in 'build_stochastic_mdp_nn' and
113 % 'build_stochastic_mdp_li'
114 switch mdp_approach
115     case 'nearest_neighbour'
116         % for each state-action pair run generate multiple samples
117         num_samples = 50;
118
119         % build stochastic MDP based on nearest neighbour
120         [T, R] = build_stochastic_mdp_nn(world, T, R, num_samples);
121
122     case 'linear_interp'
123         % build stochastic MDP based on linear interpolation
124         [T, R] = build_stochastic_mdp_li(world, T, R);
125 end
126 % =====
```

main_p2_mc_rl.m

- Load model



- Solve MDP



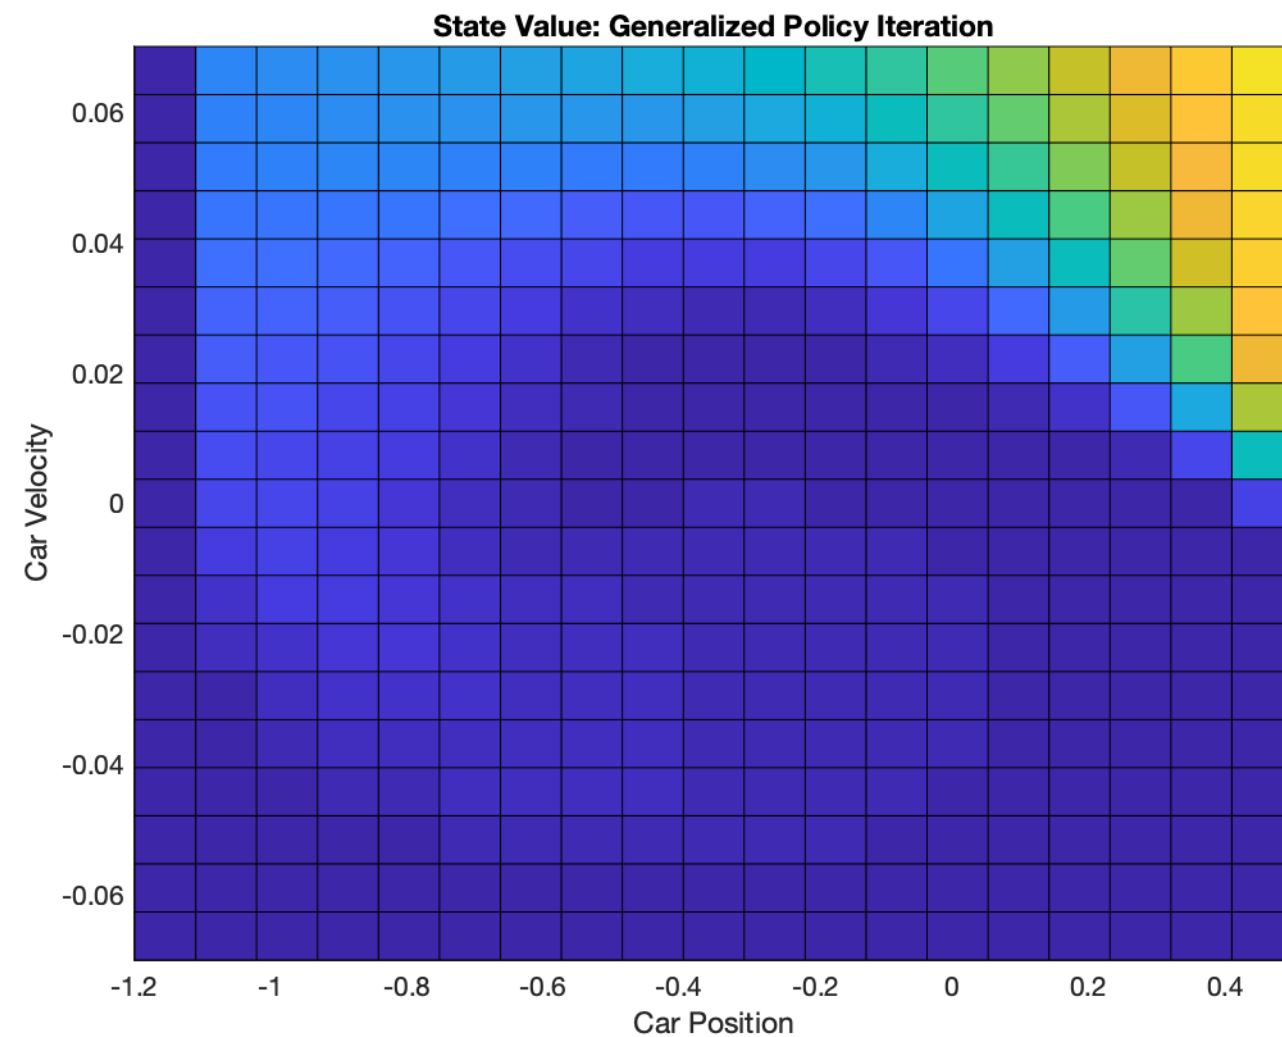
- Visualization



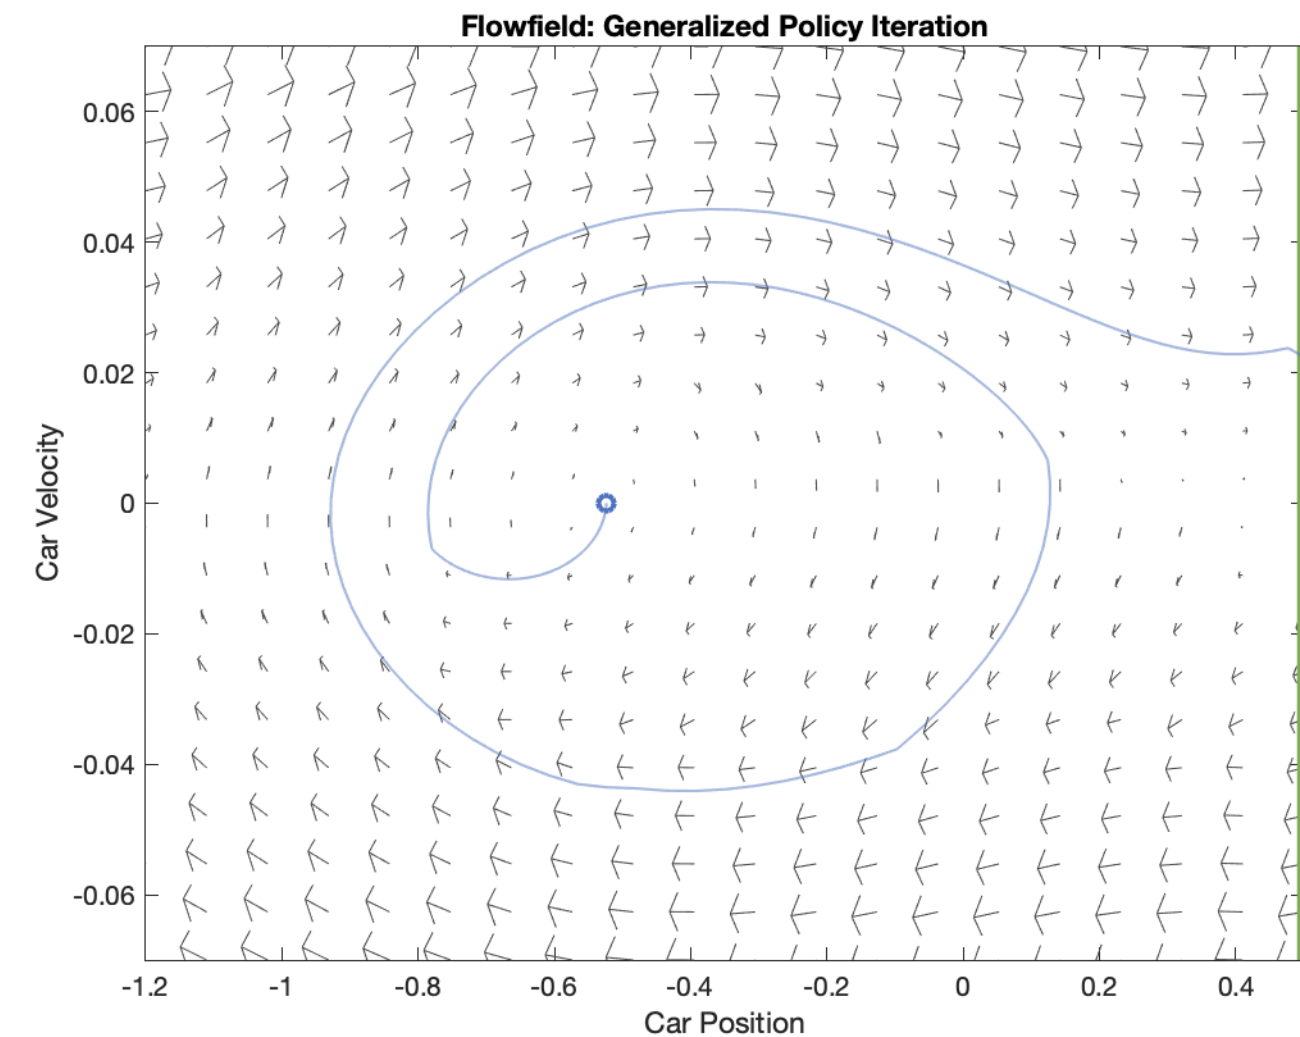
```
EDITOR PUBLISH VIEW
create_mountain_car.m x main_p2_mc_rl.m x +
35
36 %% Problem 2.2 (a)-(b) Create stochastic MDPs for the mountain car problem
37 % [TODO] Load mountain car model
38 % change model name correspondingly:
39 % (a) 'mountain_car_nn' for the nearest neighbour method
40 % (b) 'mountain_car_li' for the linear interpolation approach
41 load('./mountain_car_model/mountain_car_nn');
42
43 %% Generalized policy iteration
44 % Algorithm parameters
45 precision_pi = 0.1;
46 precision_pe = 0.01;
47 max_ite_pi = 100;
48 max_ite_pe = 100;
49
50 % Solve MDP
51 [v_gpi, policy_gpi] = generalized_policy_iteration(world, precision_pi, ...
52     precision_pe, max_ite_pi, max_ite_pe);
53
54 % Visualization
55 plot_value = true;
56 plot_flowfield = true;
57 plot_visualize = true;
58 plot_title = 'Generalized Policy Iteration';
59 hdl_gpi = visualize_mc_solution(world, v_gpi, policy_gpi, plot_value, ...
60     plot_flowfield, plot_visualize, plot_title, save_dir);
61
```


Problem 1.2 Mountain Car | Example Result with RL

- Solution with GPI




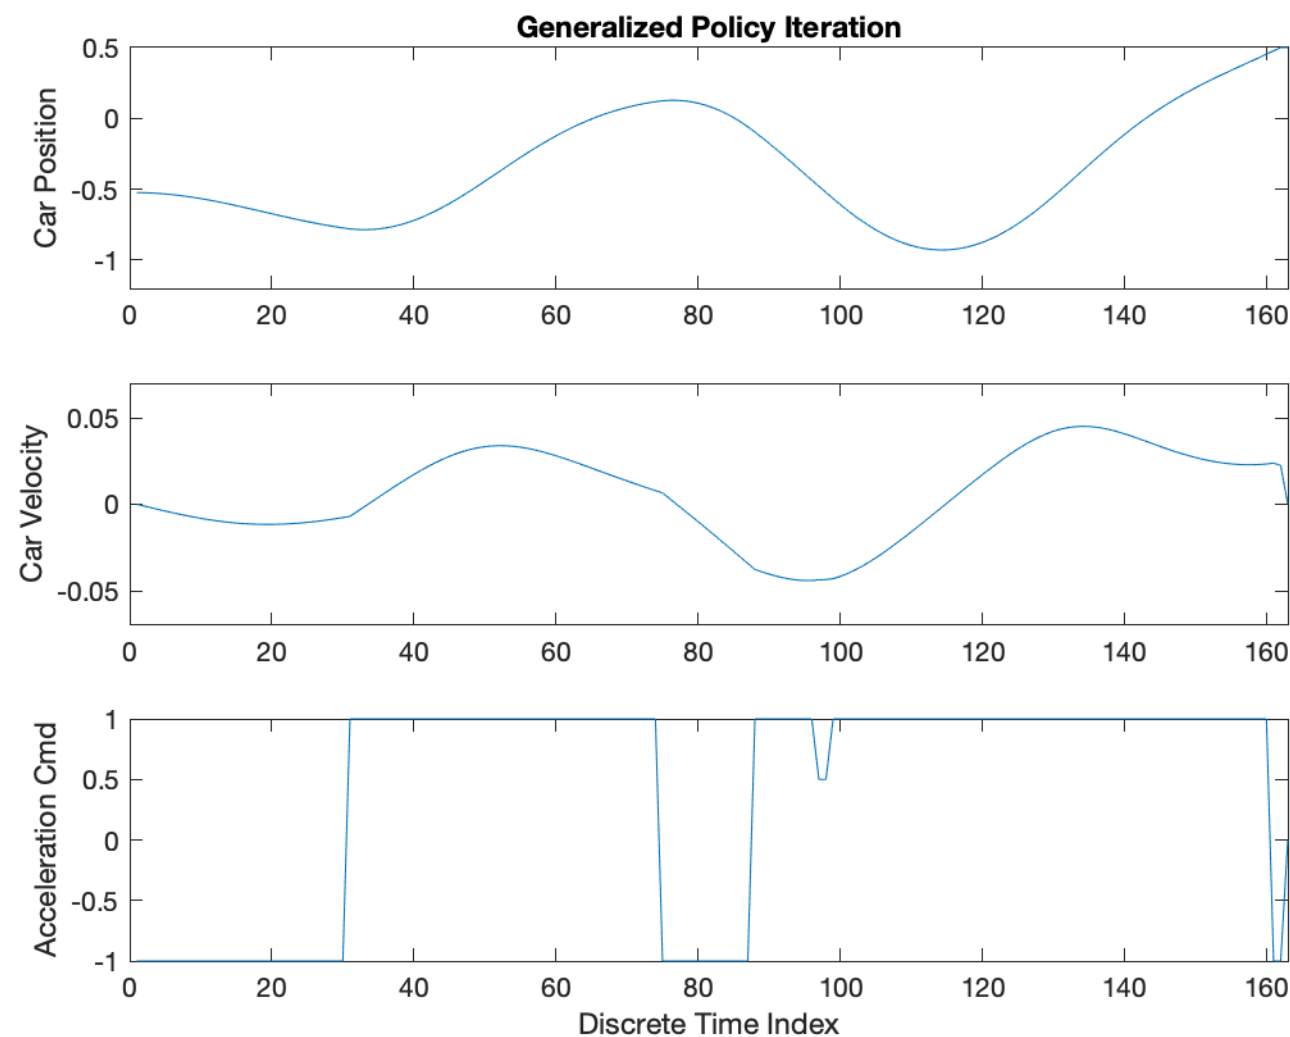
State value function (or, cost-to-go)



Resulting flowfield from the optimal policy

Problem 1.2 Mountain Car | Example Result with RL

- Solution with GPI



Link to animation: https://drive.google.com/file/d/1OxFt75O-OB8dfkJXxCmuutl0iIVLPBK_/view

Problem 1.2 Mountain Car | Code Structure (MPC Design)

main_p2_mc_mpc.m

- Nonlinear MPC
(first time step)




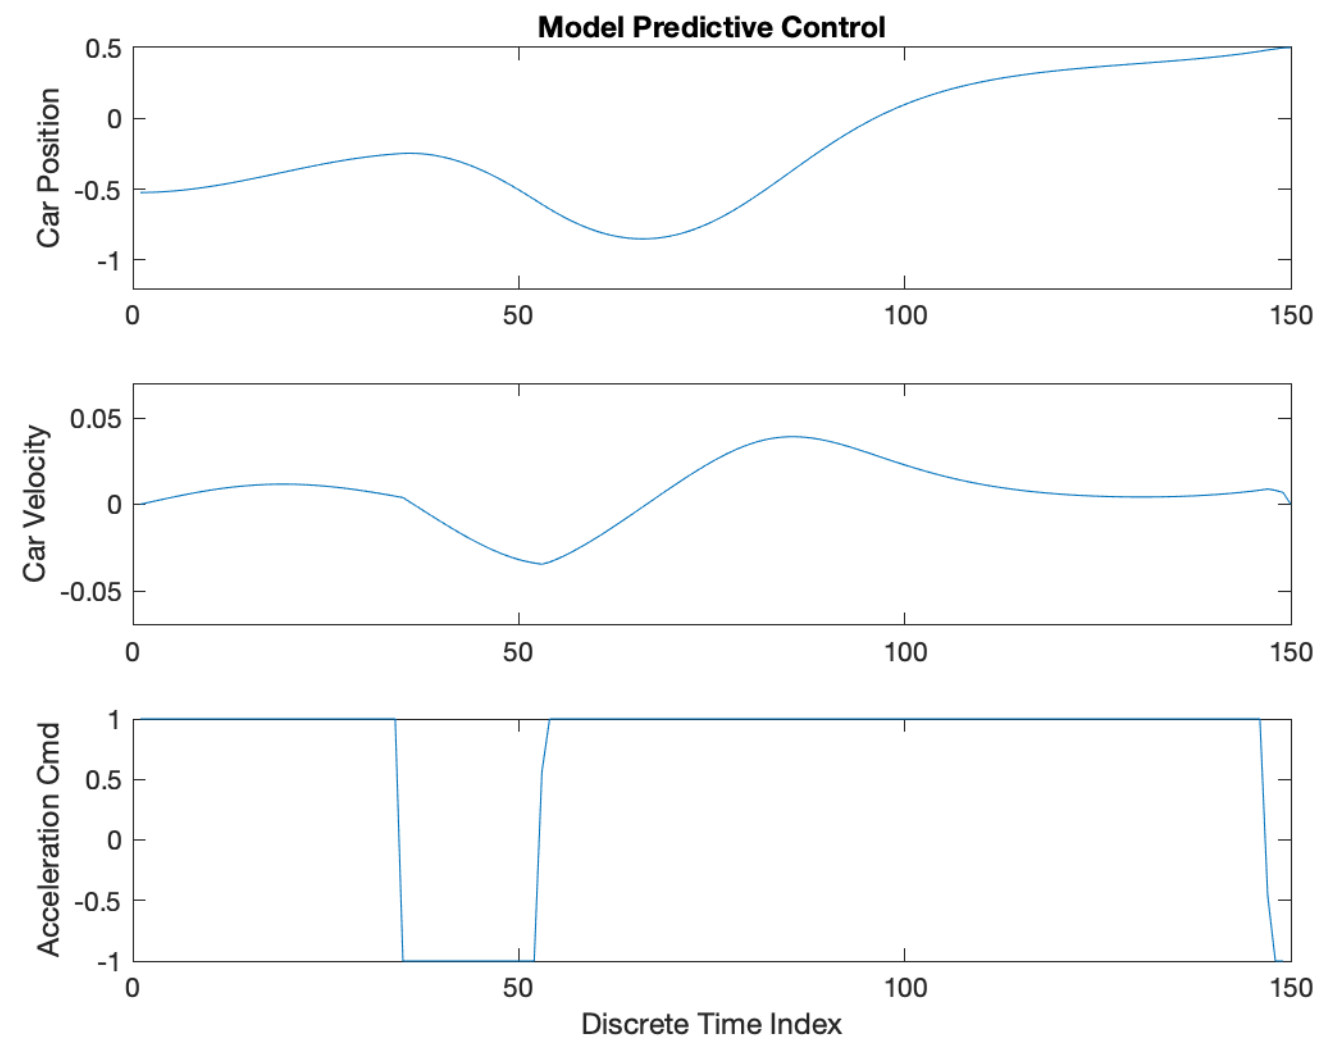
- Linearized Problem (SQP)
(subsequent time steps)



```
EDITOR PUBLISH VIEW
main_p2_mc_mpc.m
101 initial_guess = x;
102 end
103
104 % Cost function
105 sub_states = [repmat(0,n_lookahead,1); ...
106               repmat(goal_state, n_lookahead,1)];
107 fun = @(x) (x - sub_states)'*S*(x - sub_states);
108
109 % Temporary variables used in 'dyncons'
110 save('params', 'n_lookahead', 'dim_state', 'dim_action');
111 save('cur_state', 'cur_state');
112
113 % Solve nonlinear MPC
114 % x is a vector containing the inputs and states over the
115 % horizon [input,..., input, state', ..., state']^T
116 options = optimoptions(@fmincon, 'MaxFunctionEvaluations', ...
117                             1e5, 'MaxIterations', 1e5, 'Display', 'iter');
118 [x,fval] = fmincon(fun, initial_guess, [], [], [], [], ...
119                  lb, ub, @dyncons, options);
120 else
121 % ===== [TODO] QP Implementation =====
122 % Problem 2.2: (d) Quadratic Program optimizing state and
123 % action over prediction horizon
124
125 % Feedback state used in MPC updates
126 % 'cur_state' or 'cur_state_noisy'
127 cur_state_mpc_update = cur_state;
128
129 % Solve QP (e.g., using Matlab's quadprog function)
130 % Note 1: x is a vector containing the inputs and states over
131 %         the horizon [input,..., input, state', ..., state']^T
132 % Note 2: The function 'get_lin_matrices' computes the
133 %         Jacobians (A, B) evaluated at an operation point
134
135 % x = ...;
136 % =====
137
138 end
```

Problem 1.2 Mountain Car | Example Result with MPC

- Solution with MPC



Link to animation: <https://drive.google.com/file/d/192yo6gTwrGZHf9pPFfXgQbuOFvVHESkC/view>